

## Housing Price Prediction Using Linear Regression

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [ ]: data = pd.read_csv('Housingprice.csv')
```

```
In [ ]: data.head(5)
```

```
Out[ ]:   Unnamed:
0      price  lotsize  bedrooms  bathrms  stories  driveway  recroom  fullb
```

0	1	42000.0	5850	3	1	2	yes	no
1	2	38500.0	4000	2	1	1	yes	no
2	3	49500.0	3060	3	1	1	yes	no
3	4	60500.0	6650	3	1	2	yes	yes
4	5	61000.0	6360	2	1	1	yes	no

```
In [ ]: # Dropping the 'Unnamed: 0' column as it's just identifier
data.drop('Unnamed: 0', axis=1, inplace=True)
data.head(5)
```

```
Out[ ]:   price  lotsize  bedrooms  bathrms  stories  driveway  recroom  fullbase  gashw
```

0	42000.0	5850	3	1	2	yes	no	yes	no
1	38500.0	4000	2	1	1	yes	no	no	no
2	49500.0	3060	3	1	1	yes	no	no	no
3	60500.0	6650	3	1	2	yes	yes	no	no
4	61000.0	6360	2	1	1	yes	no	no	no

```
In [ ]: print('Shape of DataFrame: ', data.shape, '\n')
```

Shape of DataFrame: (546, 12)

About Dataset There are 546 rows and 12 columns in the dataset, each are:

- price: The price of the property.
- lotsize: The size of the lot in square feet.
- bedrooms: The number of bedrooms in the property.
- bathrms: The number of bathrooms in the property.
- stories: The number of stories in the property.

- driveway: Whether the property has a driveway (yes/no).
- recroom: Whether the property has a recreational room (yes/no).
- fullbase: Whether the property has a full basement (yes/no). gashw: Whether the property has gas hot water heating (yes/no).
- airco: Whether the property has central air conditioning (yes/no).
- garagepl: The number of garage places.
- prefarea: Whether the property is in a preferred location (yes/no).

```
In [ ]: data.isnull().sum()
```

```
Out[ ]: price      0
        lotsize    0
        bedrooms  0
        bathrms    0
        stories    0
        driveway   0
        recroom     0
        fullbase    0
        gashw       0
        airco       0
        garagepl    0
        prefarea    0
        dtype: int64
```

```
In [ ]: data.duplicated().sum()
```

```
Out[ ]: 1
```

```
In [ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 546 entries, 0 to 545
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   price       546 non-null   float64
1   lotsize     546 non-null   int64
2   bedrooms    546 non-null   int64
3   bathrms     546 non-null   int64
4   stories     546 non-null   int64
5   driveway    546 non-null   object
6   recroom     546 non-null   object
7   fullbase    546 non-null   object
8   gashw       546 non-null   object
9   airco       546 non-null   object
10  garagepl    546 non-null   int64
11  prefarea    546 non-null   object
dtypes: float64(1), int64(5), object(6)
memory usage: 51.3+ KB
```

```
In [ ]: data.describe()
```

	price	lotsize	bedrooms	bathrms	stories	garage
<b>count</b>	546.000000	546.000000	546.000000	546.000000	546.000000	546.000000
<b>mean</b>	68121.597070	5150.265568	2.965201	1.285714	1.807692	0.692308
<b>std</b>	26702.670926	2168.158725	0.737388	0.502158	0.868203	0.861308
<b>min</b>	25000.000000	1650.000000	1.000000	1.000000	1.000000	0.000000
<b>25%</b>	49125.000000	3600.000000	2.000000	1.000000	1.000000	0.000000
<b>50%</b>	62000.000000	4600.000000	3.000000	1.000000	2.000000	0.000000
<b>75%</b>	82000.000000	6360.000000	3.000000	2.000000	2.000000	1.000000
<b>max</b>	190000.000000	16200.000000	6.000000	4.000000	4.000000	3.000000

The dataset presents various features related to housing attributes:

- Price: Mean price stands at approximately 68, 121, *ranging from* 25,000 to \$190,000.
- Lotsize: The average lot size is about 5150 square feet, with values ranging from 1650 to 16,200 square feet.
- Bedrooms: On average, houses feature around 3 bedrooms, varying from 1 to 6.
- Bathrooms: The dataset records an average of 1.29 bathrooms per house, with a range of 1 to 4.
- Stories: The average number of stories per house is approximately 1.81, with a minimum of 1 and a maximum of 4.
- Garagepl: Houses typically offer around 0.69 garage places on average, ranging from 0 to 3.

```
In [ ]: # Converting categorical columns to numerical
categorical_columns = ['driveway', 'recroom', 'fullbase', 'gashw', 'airco',
data[categorical_columns] = (data[categorical_columns] == 'yes').astype(int)
data
```

```
Out [ ]:
```

	price	lotsize	bedrooms	bathrms	stories	driveway	recroom	fullbase	gas
0	42000.0	5850	3	1	2	1	0	1	
1	38500.0	4000	2	1	1	1	0	0	
2	49500.0	3060	3	1	1	1	0	0	
3	60500.0	6650	3	1	2	1	1	0	
4	61000.0	6360	2	1	1	1	0	0	
...	...	...	...	...	...	...	...	...	...
541	91500.0	4800	3	2	4	1	1	0	
542	94000.0	6000	3	2	4	1	0	0	
543	103000.0	6000	3	2	4	1	1	0	
544	105000.0	6000	3	2	2	1	1	0	
545	105000.0	6000	3	1	2	1	0	0	

546 rows × 12 columns

```
In [ ]: # Let's see the top 5 houses with the highest prices
highest_prices = pd.DataFrame(data.nlargest(5, ['price']))
highest_prices
```

```
Out [ ]:
```

	price	lotsize	bedrooms	bathrms	stories	driveway	recroom	fullbase	gas
161	130000.0	6000	4	1	2	1	0	1	
360	130000.0	6600	4	2	2	1	1	1	
93	128000.0	8500	3	2	4	1	0	0	
129	127000.0	4600	3	2	2	1	1	0	
374	126500.0	6420	3	2	2	1	0	0	


```
In [ ]: # Also, the top 5 houses with the lowest prices

lowest_sales = pd.DataFrame(data.nsmallest(5, ['price']))
lowest_sales
```

```
Out [ ]:
```

	price	lotsize	bedrooms	bathrms	stories	driveway	recroom	fullbase	gasl
55	25000.0	3620	2	1	1	1	0	0	
162	25000.0	2910	3	1	1	0	0	0	
232	25000.0	3850	3	1	2	1	0	0	
56	25245.0	2400	3	1	1	0	0	0	
238	26000.0	3000	2	1	1	1	0	1	

```
In [ ]: # Plotting the boxplot
plt.figure(figsize=(20, 5))
sns.boxplot(data=data, orient='h', palette='Set2')
plt.title('Boxplot of House Metrics')
plt.xlabel('Values')
plt.show()
```



```
In [ ]: # Calculate outlier bounds for 'price' and 'lotsize'
outlier_bounds_price = outlier_bounds(data['price'])
outlier_bounds_lotsize = outlier_bounds(data['lotsize'])

# Check if the outlier bounds are not None before unpacking
if outlier_bounds_price is not None and outlier_bounds_lotsize is not None:
    price_lower, price_upper = outlier_bounds_price
    lotsize_lower, lotsize_upper = outlier_bounds_lotsize

    # Create a mask for rows without outliers
    mask_no_outliers = ((data['price'] >= price_lower) & (data['price'] <= price_upper) &
                        (data['lotsize'] >= lotsize_lower) & (data['lotsize'] <= lotsize_upper))

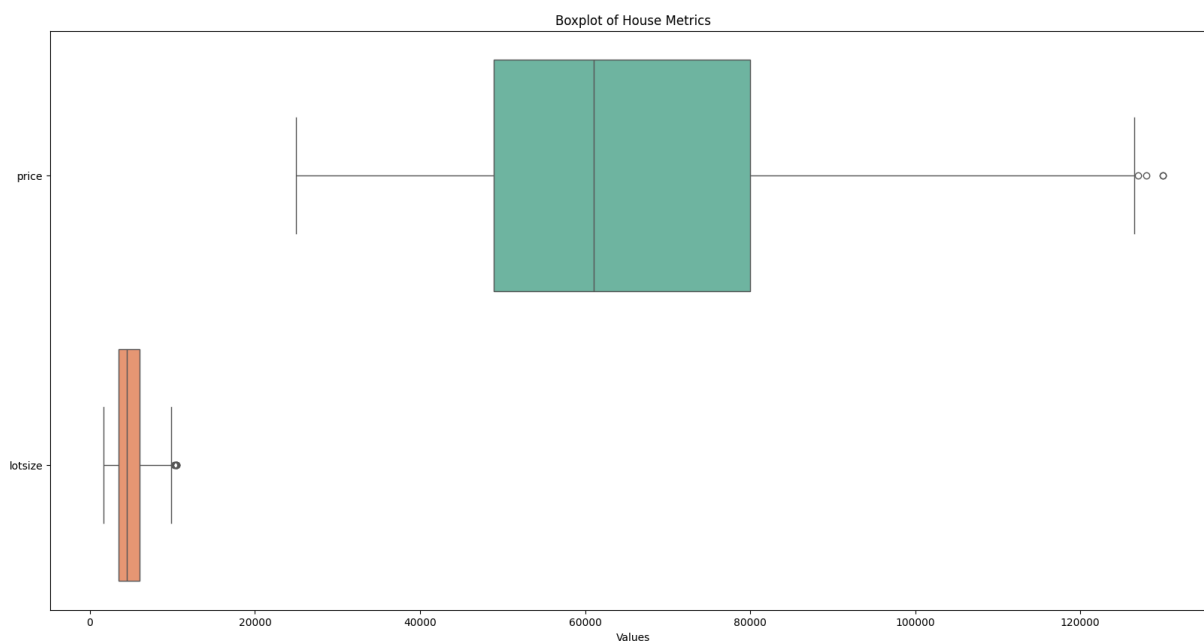
    # Apply the mask to the DataFrame
    data = data[mask_no_outliers]

else:
    print("No outliers detected.")
```

```
In [ ]: # Defining the features and target variable
X = data[['lotsize', 'bedrooms', 'bathrms', 'stories', 'driveway', 'recroom']]
y = data['price']
```


```
In [ ]: # Plot the boxplot
plt.figure(figsize=(20, 10))
```

```
sns.boxplot(data=data[['price', 'lotsize']], orient='h', palette='Set2')
plt.title('Boxplot of House Metrics')
plt.xlabel('Values')
plt.show()
```




```
In [ ]: # Removing Outliers further
data = data[(data['price'] < 80000) &
            (data['lotsize'] < 6000)]
```

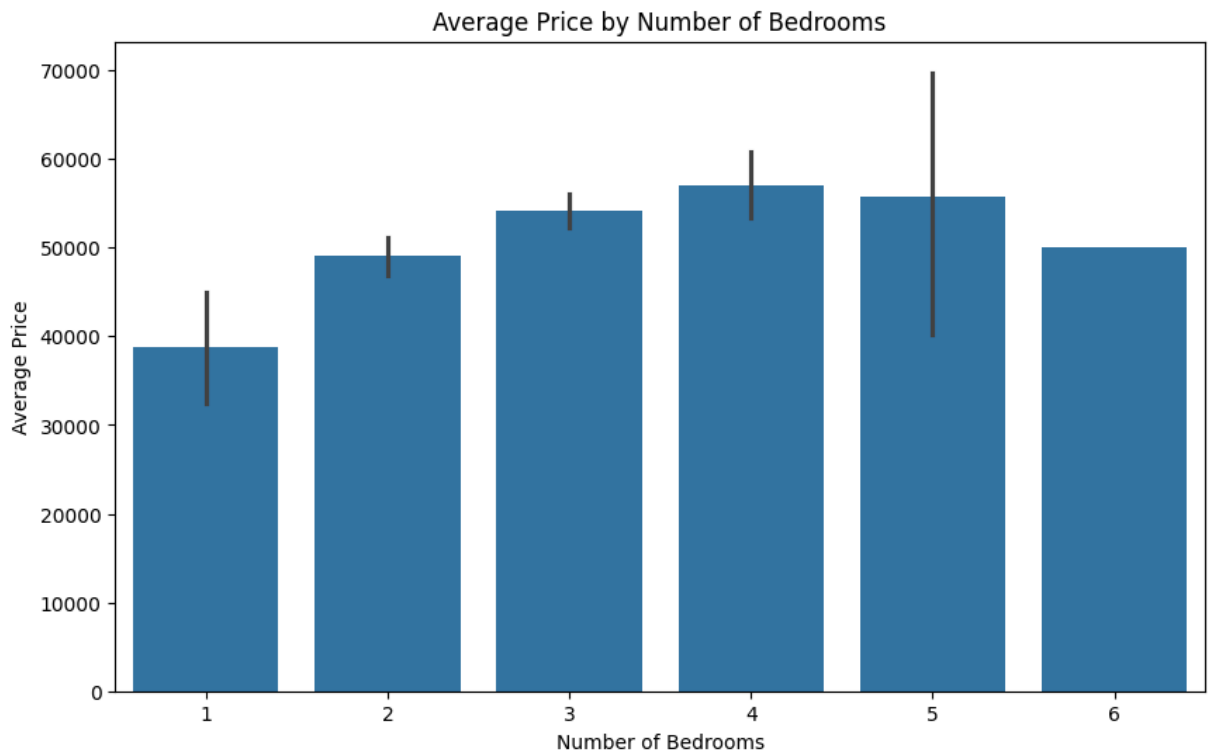
```
In [ ]: # Let's see outliers are removed or not
plt.figure(figsize=(20, 10))
sns.boxplot(data=data[['price', 'lotsize']], orient='h', palette='Set2')
plt.title('Boxplot of House Metrics')
plt.xlabel('Values')
plt.show()
```



```
In [ ]: # Barplot to visualize the distribution of houses by number of bedrooms
plt.figure(figsize=(10, 6))
sns.countplot(x='bedrooms', data=data)
plt.title('Count of Houses by Number of Bedrooms')
plt.xlabel('Number of Bedrooms')
plt.ylabel('Count')
plt.show()
```



```
In [ ]: #Barplot to visualize Average prices by Number of Bedrooms
plt.figure(figsize=(10, 6))
sns.barplot(x='bedrooms', y='price', data=data)
plt.title('Average Price by Number of Bedrooms')
plt.xlabel('Number of Bedrooms')
plt.ylabel('Average Price')
plt.show()
```

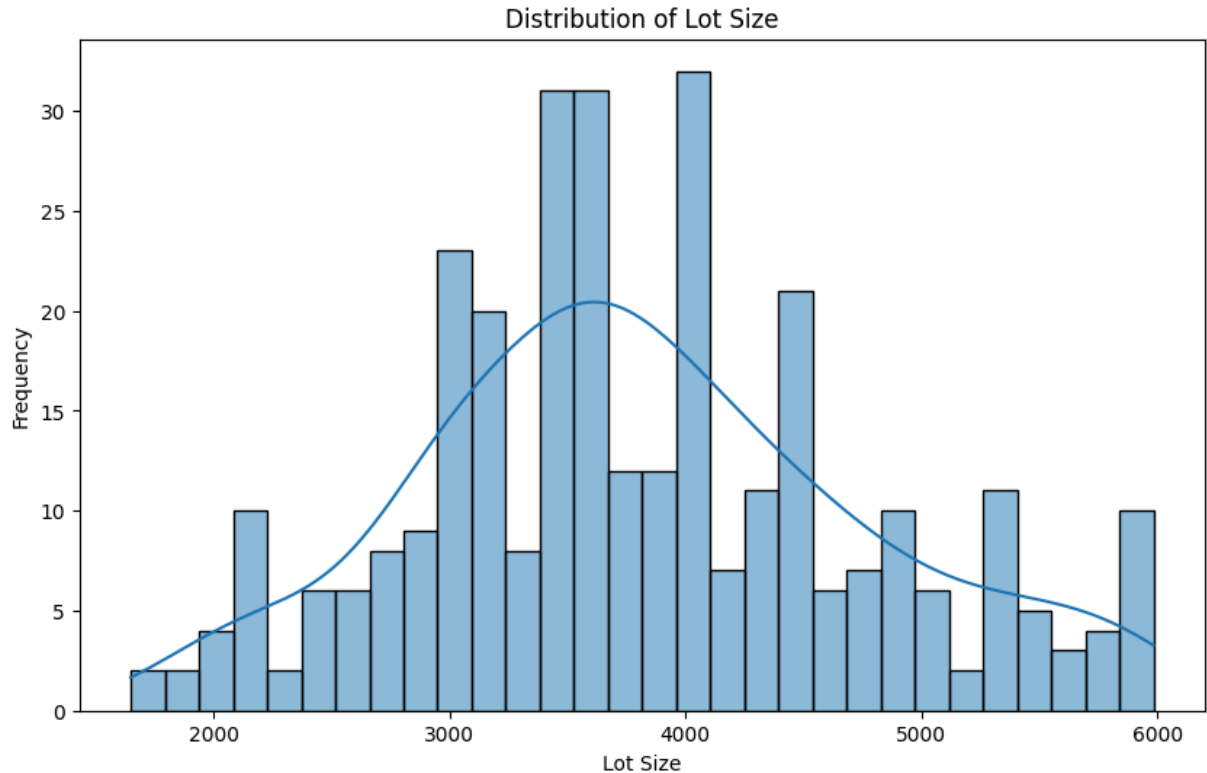


```
In [ ]: #Scatterplot to visualize relationship between Lot Size and Price
plt.figure(figsize=(10, 6))
sns.scatterplot(x='lotsize', y='price', data=data)
plt.title('Lot Size vs Price')
plt.xlabel('Lot Size')
plt.ylabel('Price')
plt.show()
```






```
In [ ]: #Histogram to visualize the distribution of lot size
plt.figure(figsize=(10, 6))
sns.histplot(data['lotsize'], bins=30, kde=True)
plt.title('Distribution of Lot Size')
plt.xlabel('Lot Size')
plt.ylabel('Frequency')
plt.show()
```



```
In [ ]: # Heatmap to visualize the correlation between variables
correlation_matrix = data.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```



## Model Building

```
In [ ]: #importing necessary libraries
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from sklearn.preprocessing import StandardScaler
```

```
In [ ]: # Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [ ]: # Standardizing the features using Standard Scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [ ]: # Modeling
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
In [ ]: # Evaluating the model
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```

```
print("Root Mean Squared Error (RMSE):", rmse)  
print("R-squared (R^2) Score:", r2)
```

Root Mean Squared Error (RMSE): 11391.130147172003  
R-squared (R^2) Score: 0.7471154326157805

```
In [ ]: accuracy_score = r2_score(y_test, y_pred)  
print("Accuracy Score:", accuracy_score)
```

Accuracy Score: 0.7471154326157805